



High Performance
Computing &
Big Data Services



Channels and Pipes with OpenCL

Emmanuel Kieffer



EUMaster4HPC




UNIVERSITÉ DU
LUXEMBOURG

Why using Channels/Pipes ?




Traditional communication scheme




- All communications go through global memory
- Can quickly become a bottleneck

Channels



- Kernel-to-kernel communication
 - Io-to-kernel communication
 - Host to kernel communication avoiding global memory
- } Independent from host

FIFO memory



- FIFO memory allows efficient communication
- Bypass Global Memory
- Streaming data application
- Channels/Pipe
 - FIFO-like communication
 - Call site is unidirectional (no default duplex communication)
 - Allows BSP-specific I/O communication with kernel compute units

Advantages



- Internal bandwidth of the FPGA hardware
- No bottleneck of using off-chip memory
- Less latency with concurrent kernel execution
- Data is consumed as it is produced which reduced storage requirements
- With Channel/Pipe communication, the hosts can launch kernels in parallel providing Performance Gains

IO Channels



- Classic communication scheme
 - Data need to be written to global memory before kernel get access to it
 - The bandwidth is limited by the PCIe bandwidth and the memory throughput
- With IO channels
 - Kernels send directly data across network interface
 - System is running at speed of network interface

How to use Channels ?



Channels creation



- Channels are specific to Intel FPGA
- Should be declared in OpenCL files (*.cl)
- Enable Intel FPGA extension of channels : `#pragma OPENCL EXTENSION cl_intel_channels : enable`
- Any built-in OpenCL or user defined type are supported
 - Structs, char, uchar, short, int, uint, long, ulong, float, vector data types
 - Type must be 1024 bits or less
- Channel FIFO depth can be defined

```
channel long a; // unbuffered channel a of type `long`  
channel long b __attribute__((depth(8))); // buffered channel b;  
channel float4 c[2]; // Channels made up of 2 float4 channels, c[0] and c[1]
```

Channels reads and write



- Read and writes data from channels are blocking functions
- Each write add a single piece of data to the channel
 - `void write_channel_intel(channel <type> channel_id, const <type> data);`
 - Ex: `write_channel_intel(channel a, (float4) vector);`
- Each read remove a single piece of data from the channel
 - `<type> read_channel_intel(channel <type> channel_id);`
 - Ex: `read_channel_intel(channel b);`
- Write blocks when the channel is fully
- Read blocks when the channel is empty

Non-blocking versions



```
bool write_channel_nb_intel(channel <type> channel_id, const <type> data);  
<type> read_channel_nb_intel(channel <type> channel_id, (float4) vector);
```

- Same functions except that the functions do not block and the pipeline is not stalled
- The bool value indicates if the operation has been successful
- Non-Blocking should be privileged when it is not guaranteed that the operation is carried out
 - Ex: I/O channels

Kernel concurrency



- To take advantage of kernel concurrency and execute kernels in parallel
- You need to create a separate command queue for each kernel

```
#define NUM_KERNELS  
...  
// Ex with C++ and std::vector  
std::vector<cl::Kernel> kernels;  
std::vector<cl::CommandQueue> queues;
```

Buffered channels



- By default, channels are unbuffered, i.e., `__attribute__((depth(20)))`
- The depth attribute specifies a **minimum** depth for the channel
- Buffered channels should be considered when the number of reads and writes are not symmetric
- Imbalance between reads/writes can be caused by conditional communication between kernels

I/O channels



- Channels providing data to the board or exporting data outside the board
 - E.g. network interfaces, PCIe interfaces, etc ...
- Board supplier has provided a BSP (board_spec.xml) which contains the I/O channels definitions

```
<channels>
  <interface name="udp_0" port="udp0_out" type="streamsource" width="256"chan_id="eth0_in"/>
  <interface name="udp_0" port="udp0_in" type="streamsink" width="256" chan_id="eth0_out"/>
</channels>
```

- I/O channel declaration using the io attribute. Reads and writes usage is the same as other channels

```
channel QUdpWord udp_in_IO __attribute__((io("eth0_in")));
channel QUdpWord udp_ou_IO __attribute__((io("eth0_out")));
```

Channels ordering



- No channel ordering by default
 - Channel calls can be executed out of order
 - Channel calls can be executed in parallel
- Risk of **DEADLOCKS** ...
- Use the `mem_fence` function to block
 - **CLK_CHANNEL_MEM_FENCE**
 - Enforce ordering

```
__kernel void producer(...) {  
    for(...) {  
        write_channel_intel(c0, ...);  
        mem_fence(CLK_CHANNEL_MEM_FENCE);  
        write_channel_intel(c1, ...);  
    }  
}  
  
__kernel void consumer(...) {  
    for(...) {  
        val = read_channel_intel(c0);  
        mem_fence(CLK_CHANNEL_MEM_FENCE);  
        val2 = read_channel_intel(c1);  
    }  
}
```

How to use Pipes ?



Pipes



- Channels are specific to Intel FPGA
 - Implemented before Khronos defined pipes for OpenCL
- Pipes should be preferred to conform with other SDK
- Intel implements pipes as a wrapper around channels
 - Channels are statically inferred, i.e., can be modified at runtime
 - All that applied to channels are also true for pipes
- AOC does not support the entire pipe specification
 - Not fully OpenCL conformant

Pipes creation

- Pipes are specified as kernel arguments with the keyword pipe
 - `read_only` or `write_only` qualifiers should be specified
- Read and Write to the pipe with `read_pipe()` and `write_pipe()` calls

```
__kernel void producer(write_only pipe uint p0){
    for(...)
        error = write_pipe(p0, &data);
}

__kernel void consumer(read_only pipe uint p0){
    for(...)
        error = read_pipe(p0, &value);
}
```

Pipes creation -- host side

- Unlike channels, pipes requires host code to work
- Similarly to `clCreateBuffer`, the `clCreatePipe` function returns a `cl_mem` object representing the pipe object
- `clSetKernelArg` has to be used to map pipe to the correct read and write kernel args to be conform with

OpenCL standard

```
cl_mem clCreatePipe(cl_context context, cl_mem_flags flags,
```

Pipe Width

```
cl_uint pipe_packet_size, cl_uint pipe_max_packets,
```

```
const cl_pipe_properties * properties,
```

```
cl_int *errcode_ret)
```

Pipe Depth

Pipe Properties

Pipes attributes

- Contrary to channels, pipes are non-blocking by default
- Add `__attribute__((blocking))` for blocking behavior

```
__kernel void producer(write_only pipe uint __attribute__((blocking)) p0);  
__kernel void consumer(read_only pipe uint __attribute__((blocking)) p0);
```

- The depth attribute can be used to specify the minimum depth of the pipe

```
#define SIZE 100  
__kernel void producer(write_only pipe uint __attribute__((depth(SIZE))) p0);  
__kernel void consumer(read_only pipe uint __attribute__((depth(SIZE))) p0);
```

- I/O pipes use the io attributes, e.g., `__attribute__((io("eth0_in")))`

Difference with Standard OpenCL pipes




- Standard OpenCL pipes are created at execution when the host calls `clCreatePipe`
 - For Intel, the pipes are created before the host code is executed
 - Intel pipes implementation is static
 - The standard expect a dynamic implementation
- Standard pipes do not require the same `pipe_id` on the read and the write sides
 - The standard relies on the `clSetKernelArg` function call to know to match read side and write side
 - The Intel pipes need the same `pipe_id` to match read side and write side

Using channels and pipes



Channels/Pipes behavior

- Written data into a channel or a pipe remain valid during the kernel program lifetime on FPGA
- What is produced should be consumed or rejected by the consumer
- Ex: 6 elements produced and transmitted to consumer. The consumer read 3 elements at a time
 - There can be 2 read call for each write



Behavior with NDRange kernels



- Work-items have no specific ordering in NDRange kernels
- AOC enforces a work-item ordering with channels.
- These are the rules:
 - Pipelined architecture allows only reads and write once per clock cycle across different work-items
 - Smaller work-item IDs are executed first
 - “Threads” proceed in work-item and work-group order (dimX -> dimY -> dimZ)
 - Thread dependent workflow should be avoided
- For single-work item, consistency and in-order execution are ensured

Global restrictions



- Channels can have multiple read call sites but only a single call site
- Pipes are more restricted. Only a single read call and single write call site
- No loop unrolling when loops contain channels or pipes
- Kernels with channels cannot be vectorized → no `num_simd_work_items`
- Kernels with channels cannot be replicated → no `num_compute_unit`
- Static indexing mandatory for arrays of channels because AOC needs the information at compile time

Best practices

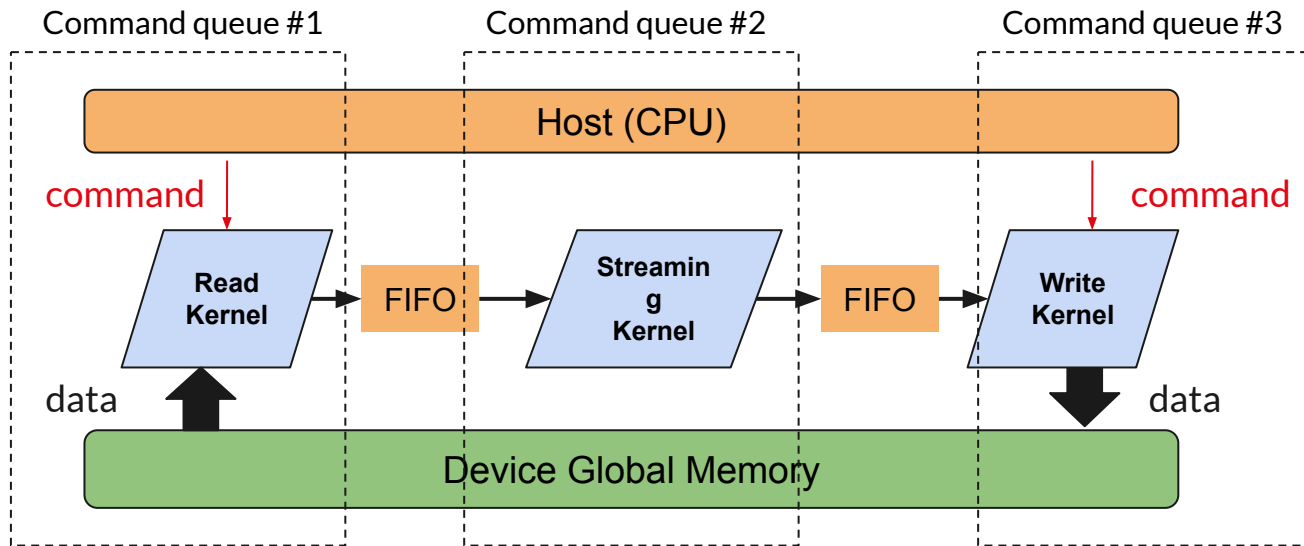
- Prefer single work-item kernels over NDRange kernels
- Plan channels/pipes connections by decomposing a single large kernels into multiple small kernels
- Keep the number of channels reasonable → aggregate data when possible else use different channels
- Don't use non-blocking versions and loop to wait data
 - Replace it with blocking channels/pipes to avoid significant waste
- Channels/Pipes can be emulated
 - Compilation for emulation, default channel depth is different from the default channel depth generated when your kernel is compiled for hardware. See [Emulating Channel Depth](#) to change this behavior
- I/O channels emulation more problematic. Refer to [Emulating Applications with a Channel That Reads or Writes to an I/O Channel](#) for more details

Channels vs Pipes



- Pipes are partially conformant with OpenCL standard
 - They can be used for host pipes, i.e. bypassing the FPGA global memory
- Channels can be used for autorun kernels (no input arguments needed)
 - No host code, less verbose
 - Should be preferred for FPGA implementation
- Both have same usage and performance


Channels/Pipes use case



- Read kernel transfers data from DDR to channels
- Streaming kernels read input channels, process data and write data to output channel
- Write kernel transfers data to DDR
- **3 different command queues**

Advised Design Model


- When dealing with a large amount of data -- generalize previous use case



- When dealing with a large amount of data

Workflow

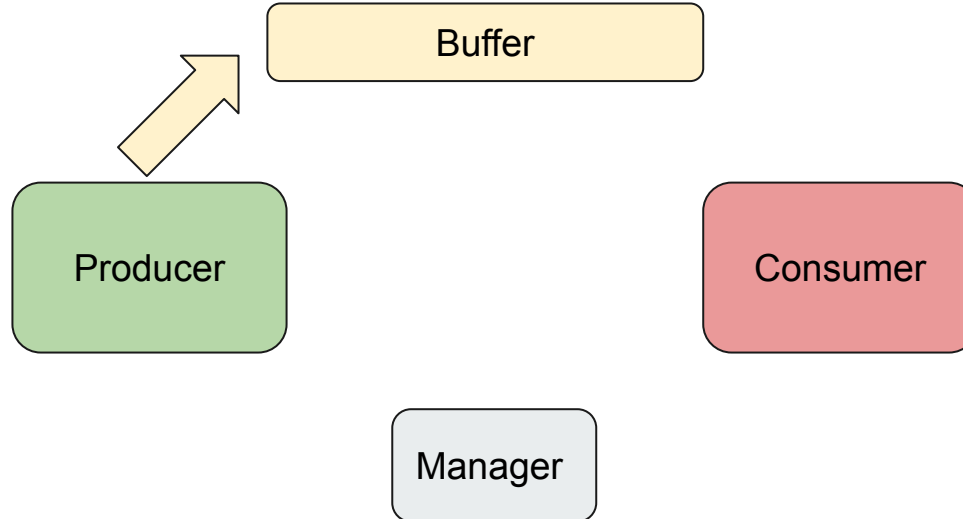
- Manager sends a token to the producer



- When dealing with a large amount of data

Workflow


- The producer write data the specific buffer region based on information inside the token



- When dealing with a large amount of data

Workflow


- The producer transmits the token to the consumer which takes over



- When dealing with a large amount of data

Workflow


- The consumer releases the buffer and send back the token to the manager



Using host pipes



Host pipes



- Host to kernel communication avoiding global memory
- Transfer data to/from the kernel without using global memory
- Only with pipes

Channels creation



- Enable Intel FPGA extension: `#pragma OPENCL EXTENSION cl_intel_fpga_host_pipe : enable`
- Use `__attribute__((intel_host_accessible))` when declaring the host pipe argument in the cl file
- Same usage than classical pipes

```
__kernel void reader(__attribute__((intel_host_accessible)) read_only pipe uint4 host_in){...}  
__kernel void writer(__attribute__((intel_host_accessible)) write_only pipe uint4 host_out){...}
```

Host pipe -- host side

- Unlike channels, pipes requires host code to work
- Create the pipe using `clCreatePipe` function
 - Use the flags `CL_MEM_HOST_READ_ONLY` or `CL_MEM_HOST_WRITE_ONLY`
- `clSetKernelArg` has to be used to map pipe to the correct read and write kernel args

```
cl_mem read_pipe = clCreatePipe(context, CL_MEM_HOST_READ_ONLY, sizeof(uint4), 64, NULL, &error);
cl_mem write_pipe = clCreatePipe(context, CL_MEM_HOST_WRITE_ONLY, sizeof(uint4), 64, NULL, &error);

clSetKernelArg(kernel, 0, sizeof(cl_mem), (void*)&read_pipe);
clSetKernelArg(kernel, 1, sizeof(cl_mem), (void*)&write_pipe);
```

Reading from and writing to host pipes



Function	Use case
<code>cl_int clReadPipeIntelFPGA(cl_mem pipe, gentype *ptr);</code>	<ul style="list-style-type: none">• Read one data from pipe• Non-Blocking. Return 0 if successful.
<code>cl_int clWritePipeIntelFPGA(cl_mem pipe, gentype *ptr);</code>	<ul style="list-style-type: none">• Read one data from pipe• Non-Blocking. Return 0 if successful.
<code>cl_int clMapHostPipeIntelFPGA(cl_mem pipe, ...);</code>	Creates buffer to provide data for reads and writes made up of multiple words
<code>cl_int clUnmapHostPipeIntelFPGA(cl_mem pipe, ...);</code>	Frees buffer to provide data for reads and writes made up of multiple words